

In [30]:

```
1 # Set up environment with correct dependencies
2 using Pkg
3 Pkg.activate(".")
4 Pkg.instantiate()
```

Activating environment at `~/GitHub/MathSys/teaching/MA934/MA934-slides/Project.toml`

In [31]:

```
1 using Plots
2 using LaTeXStrings
3 using Revise
4 using LinearAlgebra
5 using DualNumbers
6 using Random
7 using Distributions
8
9
10 pyplot()
11 # Set default fonts for all plots
12 #fnt = Plots.font("DejaVu Sans", 8.0)
13 #default(titlefont=fnt, guidefont=fnt, tickfont=fnt, legendfont=fnt)
```

Out[31]:

Plots.PyPlotBackend()

In [32]:

```

1  # Need Golden Section Search
2  include("files/code/gss.jl")
3
4  function gradient_descent_plot(f, gamma, xpos, ypos)
5      n = 100
6      x = range(-10.0, stop=10.0, length=n)
7      y = range(-10.0, stop=10.0, length=n)
8
9      xgrid = repeat(x', n, 1)
10     ygrid = repeat(y, 1, n)
11     z = zeros(n, n)
12
13     for i in 1:n
14         for j in 1:n
15             z[i:i, j:j] .= f([x[j]; y[i]], gamma)
16         end
17     end
18     p = contour(xgrid, ygrid, z, colors="black", linewidth=1.0, colorbar=false,
19     plot!(xpos, ypos, label="")
20     scatter!(xpos, ypos, label="")
21     return p
22 end
23
24 function gradient_descent_convergence_plot(xpos1, ypos1, xpos2, ypos2, labels)
25     p=plot()
26     n = length(xpos1)
27     steps = 1:n
28     dist = [sqrt(xpos1[i]^2+ypos1[i]^2) for i in steps]
29     p = plot(steps, dist, yscale=:log10, label="", color=:blue, xlabel="# itera
30     scatter!(steps, dist, color=:blue, label=labels[1])
31
32     n = length(xpos2)
33     steps = 1:n
34     dist = [sqrt(xpos2[i]^2+ypos2[i]^2) for i in steps]
35     plot!(steps, dist, yscale=:log10, label="", color=:green)
36     scatter!(steps, dist, color=:green, label=labels[2])
37
38     return p
39 end
40
41 function gradient_descent(f, df, gamma)
42     x0 = [10.0, 1.0]
43
44     g = df(x0, gamma)
45     n = 0
46     x = x0
47     dx = - g / norm(g)
48
49     steps = Int[]
50     xposition = Float64[]
51     yposition = Float64[]
52
53     while n<20 && norm(g)>1E-5
54         append!(steps, n); append!(xposition, x[1]); append!(yposition, x[2]);
55         F(lambda) = f(x + lambda*dx, gamma)
56         lmin = gss.minimise(F, [-100.0, 100.0, 0.0], 1.0E-6)
57         x = x + lmin*dx
58         g = df(x, gamma)
59         dx = -g/norm(g)

```

```

60         n+=1
61     end
62     return xposition, yposition
63 end
64
65 function gradient_descent(f, df, gamma)
66     x0 = [10.0, 1.0]
67
68     g = df(x0, gamma)
69     n = 0
70     x = x0
71     dx = - g / norm(g)
72
73     steps = Int[]
74     xposition = Float64[]
75     yposition = Float64[]
76
77     while n<20 && norm(g)>1E-5
78         append!(steps,n); append!(xposition, x[1]); append!(yposition, x[2]);
79         F(lambda) = f(x + lambda*dx, gamma)
80         lmin = gss.minimise(F, [-100.0, 0.0, 100.0], 1.0E-6)
81         x = x + lmin*dx
82         g = df(x, gamma)
83         dx = -g/norm(g)
84         n+=1
85     end
86     return xposition, yposition
87 end
88
89 function ADplot()
90     xgrid1 = range(-1.0,stop=1.0,length=101)
91     xgrid2 = range(-1.0,stop=1.0,length=51)
92     z = [Dual(x, 1.0) for x in xgrid2]
93     y = sin.(2*pi*xgrid1)
94     yprime = 2*pi*cos.(2*pi*xgrid1)
95     yz = sin.(2*pi*z)
96     p = plot(xgrid1, y, label=L"y(x) = \sin(2\pi x)", color=:red, ylims=(-7,16))
97     plot!(xgrid1, yprime, label=L"y(x) = 2\pi\cos(2\pi x)", color=:blue)
98     scatter!(xgrid2, realpart.(yz), label="Real[y(z)]", color=:red)
99     scatter!(xgrid2, dualpart.(yz), label="Dual[y(z)]", color=:blue)
100    return p
101 end

```

WARNING: replacing module gss.

Out[32]:

ADplot (generic function with 1 method)

MA934


Nonlinear optimisation

Minimisation in 1 dimension

Given a function, $f(x)$ of a single variable, the task is to find a minimum of f .



An ordered triple of points (a, c, b) is said to *bracket* a minimum of f if $f(c) < f(a)$ and $f(c) < f(b)$.

Line search: evaluate f at a new point, x , to construct a smaller bracketing triple. Iterate until a desired accuracy is reached.



Golden section search is a way to organise this search in an optimal way.

Golden section search



Most efficient to choose x in the larger of the two subintervals:

Golden section search: : choosing new point, x

Case 1: $[a, c]$ is the larger subinterval : new bracketing triple is either (a, x, c) or (x, c, b) . Width of the new bracketing triple is independent of which outcome if :

$$c - a = b - x.$$

So we choose

$$x = a + b - c.$$

Case 2: if $[c, b]$ is the larger subinterval: new bracketing triple is either (a, c, x) or (c, x, b) . Width of the new bracketing triple is independent of which outcome if :


$$x - a = b - c.$$

So we again choose

$$x = a + b - c.$$

Golden section search : choosing initial "shape"

Denote:



The idea is to choose c such that the ratio of the width of the shorter subinterval to the width of the longer one remains constant between iterations.

- p : width of the longer subinterval in the old triple
- q : width of the shorter subinterval in the old triple.
- r width of shorter subinterval in the *new* triple.

There are several cases:


Golden section search : choosing initial "shape"

Case 1:

$$\text{Case 1(A) : } \frac{r}{p-r} = \frac{q}{p} \quad \text{Case 1(B) : } \frac{r}{q} = \frac{q}{p}$$

Case 2: (get same equations)

$$\text{Case 2(A) : } \frac{r}{q} = \frac{q}{p} \quad \text{Case 2(B) : } \frac{r}{p-r} = \frac{q}{p}$$



Eliminating r gives

$$\left(\frac{q}{p}\right)^2 + \left(\frac{q}{p}\right) - 1 = 0$$

$$\Rightarrow \left(\frac{q}{p}\right) = \frac{\sqrt{5} \pm 1}{2}.$$

Golden section search

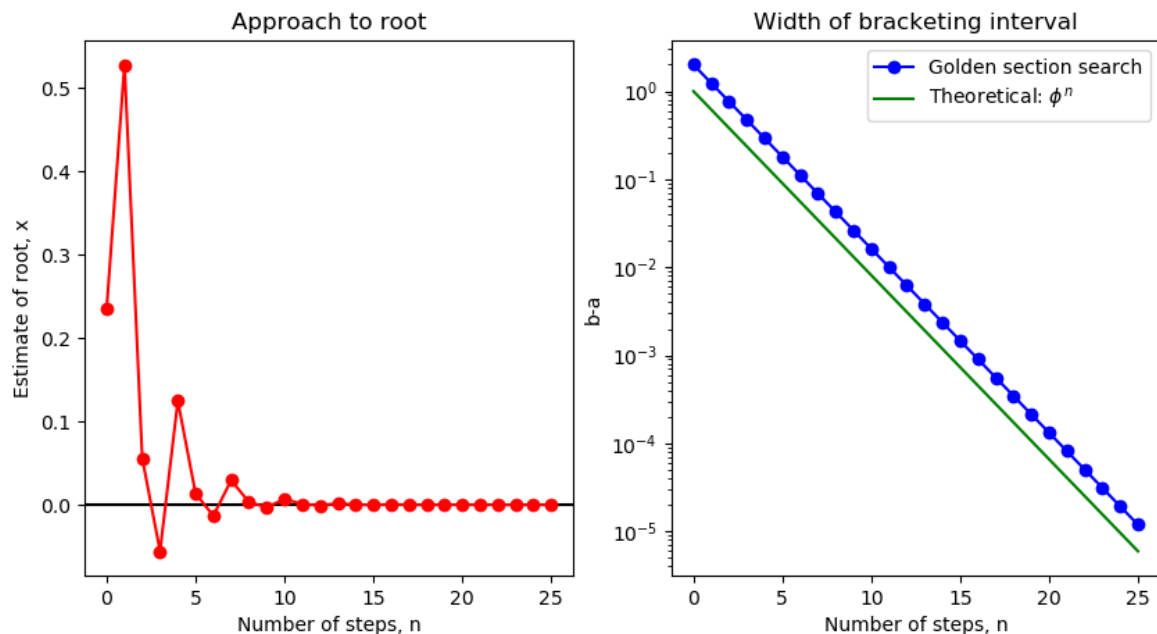
We choose the - sign since we assumed $\frac{q}{p} < 1$ ($\frac{q}{p} = \phi$, 1/golden ratio).

Since x is already determined, $x = a + b - c$, if we start with the correct ratio, $\frac{q}{p}$, this will be preserved when as we iterate.

Convergence is exponential in the number of iterations: width of interval after n iterations is $(b - a) \phi^n$.

Golden section search: convergence


Finding the minimum of $f(x) = x^2$ at $x = 0$:



Golden section search : implementation

Implementation is simplified using temporary variables x_1 and x_2 to store the points x and c (in Case 1) or c and x (in Case 2) and f_1 and f_2 to store the associated function values.

Regardless of the order in which previous points were evaluated, the new triple will centre on the point with the smallest value of f found so far. Thus by comparing f_1 to f_2 , there are only two cases:



Minimisation in \mathbb{R}^n : line minimisation

Given a function, $f : \mathbb{R}^n \rightarrow \mathbb{R}$, a position vector, $\mathbf{x} \in \mathbb{R}^n$, and a direction vector, $\mathbf{v} \in \mathbb{R}^n$, the *line minimiser* of f from \mathbf{x} in the direction \mathbf{v} is the point

$$\mathbf{x}^* = \mathbf{x} + \lambda^* \mathbf{v}$$

where

$$\lambda^* = \arg \min_{\lambda} f(\mathbf{x} + \lambda \mathbf{v}).$$

Note that although f is a function of n variables, this minimisation with respect to λ is one dimensional and can be done, for example, using Golden Section Search.

Minimisation in \mathbb{R}^n : gradient descent

Gradient descent is one a family of unconstrained optimisation algorithms that can be used when the gradient of the objective function is known or computable.

- Given $f : \mathbb{R}^n \rightarrow \mathbb{R}$ and any point, $\mathbf{x} \in \mathbb{R}^n$, $-\nabla f(\mathbf{x})$ points in the direction of steepest decrease of f at \mathbf{x} .
- Idea is to keep going "downhill" until a minimum is reached

Gradient descent algorithm

Start from a point, \mathbf{x}_0 .

1. Calculate the unit direction vector

$$\mathbf{v}_n = -\frac{\nabla f(\mathbf{x}_n)}{|\nabla f(\mathbf{x}_n)|}.$$

2. Perform line minimisation in direction of \mathbf{v}_n :

$$\lambda^* = \arg \min_{\lambda} f(\mathbf{x}_n + \lambda \mathbf{v}_n)$$

3. Move to new position $\mathbf{x}_{n+1} = \mathbf{x}_n + \lambda^* \mathbf{v}_n$.

4. Repeat until $|\nabla f(\mathbf{x}_n)| < \epsilon_{\text{tol}}$.

Normalisation of \mathbf{v}_n is not strictly necessary but helps keep accuracy as minimum is approached.

Note: any descent directions will work - gradient is the most efficient.

Gradient descent: example in \mathbb{R}^2

Consider the function

$$f(x, y) = \frac{1}{2} (x^2 + \gamma y^2).$$

The parameter $\gamma > 0$ controls the ellipticity. The gradient is

$$\nabla f(x, y) = (x, \gamma y)^T.$$

In [33]:

```
1 # Parameter to control ellipticity of
2 gamma = 10.0
3 # Define quadratic function of two variables
4 f(x, gam) = 0.5*(x[1]^2.0 + gam*x[2]^2.0)
5 # Define corresponding gradient
6 df(x, gam) = [x[1], gam*x[2]]
7
```

Out[33]:

```
df (generic function with 1
method)
```

Gradient descent: example in \mathbb{R}^2


In [34]:

```

1 gamma=4.0
2 xpos1, ypos1 = gradient_descent(f, df
3 gradient_descent_plot(f, gamma, xpos1

```

Out[34]:




In [35]:

```

1 gamma=10.0
2 xpos2, ypos2 = gradient_descent(f, df
3 gradient_descent_plot(f, gamma, xpos2

```

Out[35]:



Gradient descent: Why the zig zags?

Remember that \mathbf{x}_{n+1} is a line minimiser of f from \mathbf{x}_n in the direction of \mathbf{v}_n .

Letting $g(\lambda) = f(\mathbf{x}_n + \lambda \mathbf{v}_n)$, we must have

$$\begin{aligned} \frac{dg}{d\lambda}(\lambda^*) &= 0 \\ \Rightarrow \sum_{i=1}^n \frac{\partial f}{\partial x_i}(\mathbf{x}_n + \lambda^* \mathbf{v}_n) v_{ni} &= 0 \\ \Rightarrow \nabla f(\mathbf{x}_{n+1}) \cdot \mathbf{v}_n &= 0. \end{aligned}$$

Each step of the gradient descent algorithm is perpendicular to the previous one.

Gradient descent: convergence

- convergence is exponential but depends on γ .
- larger γ is harder - narrower valley requires more zig zags.
- this latter feature means vanilla gradient descent is rarely used in practice.


In [36]:

```

1 gradient_descent_convergence_plot(xpc

```

Out[36]:



Numerical calculation of derivatives

Explicit calculation of $\nabla f(\mathbf{x})$ is sometimes inconvenient. In such cases it is helpful to be able to calculate derivatives numerically. Several approaches:

- Finite difference approximation: the workhorse.
- Automatic differentiation: uses dual numbers
- Spectral methods: uses Fourier transforms

Dual Numbers: basic arithmetic

Dual numbers are of the form $z = x + \varepsilon y$ with $x, y \in \mathbb{R}$ and $\varepsilon^2 = 0$.

Addition rule:

$$\begin{aligned} z_1 + z_2 &= (x_1 + \varepsilon y_1) + (x_2 + \varepsilon y_2) \\ &= (x_1 + x_2) + \varepsilon (y_1 + y_2). \end{aligned}$$

Dual numbers: division and powers

Division is defined using the conjugate:

$$\begin{aligned} \frac{z_1}{z_2} &= \frac{z_1 \bar{z}_2}{z_2 \bar{z}_2} \\ &= \frac{(x_1 + \varepsilon y_1)(x_2 - \varepsilon y_2)}{x_2^2} \\ &= \frac{x_1 x_2 + \varepsilon (y_1 x_2 - x_1 y_2)}{x_2^2} \\ &= \frac{x_1}{x_2} + \varepsilon \left(\frac{y_1 x_2 - x_1 y_2}{x_2^2} \right). \end{aligned}$$

This property trivially extends to polynomial functions and *formally* extends to real analytic functions via their Taylor series:

$$\begin{aligned} f(x + \varepsilon y) &= \sum_{k=0}^{\infty} \frac{1}{k!} f^{(k)}(x) \varepsilon^k y^k \\ &= f(x) + \varepsilon y f'(x). \end{aligned}$$

Also the chain rule works (check). If $h(x) = f(g(x))$:

$$h(x + \varepsilon) = f(g(x)) + \varepsilon f'(g(x)) g'(x).$$

Multiplication rule:

$$\begin{aligned} z_1 * z_2 &= (x_1 + \varepsilon y_1) * (x_2 + \varepsilon y_2) \\ &= (x_1 x_2) + \varepsilon (x_1 y_2 + x_2 y_1). \end{aligned}$$

Dual conjugate:

$$\bar{z} = x - \varepsilon y.$$

As with the complex numbers, $z \bar{z}$ is purely real:

$$z \bar{z} = (x + \varepsilon y) * (x - \varepsilon y) = x^2.$$

Division is not defined for $\text{Re}(z_2) = 0$ so, unlike \mathbb{C} , the dual numbers do not form a field.

Powers are obtained using the binomial theorem:

$$\begin{aligned} (x + \varepsilon y)^n &= \sum_{k=0}^n \binom{n}{k} x^{n-k} (\varepsilon y)^k \\ &= \binom{n}{0} x^n + \binom{n}{1} x^{n-1} \varepsilon y \\ &= x^n + \varepsilon y n x^{n-1}. \end{aligned}$$

Dual numbers: automatic differentiation

Notice that for $f(x) = x^n$,

$$f(x + \varepsilon y) = f(x) + \varepsilon y f'(x).$$

Evaluating $f(x) = x^n$ at $x + \varepsilon$ gives the derivative of $f(x)$ as the dual component.

This is called *automatic differentiation*.

Dual numbers in Julia

Advantages of automatic differentiation:

- no need to work out complicated analytic formulae for derivatives
- `Dual[z]` correct to the same precision as `Real[z]`.

Dual arithmetic is supported in Julia by various packages, the most basic of which is `DualNumbers.jl`.

In [37]:

```

1 using DualNumbers
2 z1 = Dual(2.0, 3.0)
3 z2 = Dual(1.0, 2.0)
4 z1*z2

```

Out[37]:

2.0 + 7.0ε


In [38]:

```

1 x = range(-1.0, stop=1.0, length=101)
2 z = [Dual(a, 1.0) for a in x]
3 y = sin.(2*pi*z)
4
5 ADplot()

```

Out[38]:



Automatic differentiation for multivariate functions

In the formulation presented above, automatic differentiation of multivariate functions requires a separate function evaluation for each partial derivative.

For a multivariate function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, to calculate the partial derivative with respect to x_i , we have to add ε to the i^{th} component of the argument and then take the dual part. For example, in \mathbb{R}^2 we have

$$f(x + \varepsilon, y) = f(x, y) + \varepsilon \frac{\partial f}{\partial x}(x, y)$$

$$f(x, y + \varepsilon) = f(x, y) + \varepsilon \frac{\partial f}{\partial y}(x, y).$$

Minimisation of sums

Problems in statistical inference, ML and data science frequently produce optimisation problems that involve minimising a sum:

$$\min_{\beta} f(\beta) = \min_{\beta} \frac{1}{n} \sum_{i=1}^n L(\beta, \mathbf{x}_i).$$

Often β represents some model parameters, \mathbf{x}_i represents the i^{th} observation in the training data set and L is a loss function of some kind.

Can be solved with GD by calculating the gradient with respect to β :

$$(\nabla f(\beta))_j = \frac{1}{n} \sum_{i=1}^n \frac{\partial L}{\partial \beta_j}(\beta, \mathbf{x}_i)$$

Stochastic Gradient Descent

If the training data is large ($n \gg 1$), then calculating the gradient can become very expensive. Stochastic Gradient Descent addresses this issue.

Idea is to estimate the gradients as:

$$(\nabla f(\beta))_j \approx (\tilde{\nabla} f(\beta))_j = \frac{\partial L}{\partial \beta_j}(\beta, \tilde{\mathbf{x}}_i)$$

where $\tilde{\mathbf{x}}_i$ is the i^{th} training data point from the training data set having first undergone a random shuffle.

Stochastic Gradient Descent : learning rate

With SGD, we typically do not perform full line minimisations (due to sampling noise). Instead the basic update rule is

$$\beta_{n+1} = \beta_n - \zeta_n (\tilde{\nabla} f(\beta_n)).$$

ζ_n is a (decreasing) function of n often called the *learning rate*. For example,

$$\zeta_n = \frac{\zeta_0}{1 + \zeta_1 n}$$

ζ_0 and ζ_1 are *hyperparameters*. Note that $\sum_{n=1}^{\infty} \zeta_n$ is a divergent series. Prevents SGD from stalling.

Due to sampling noise, SGD only converges to a "noise ball" around the minimum of $f(\beta)$ rather than to the true minimum.

In [39]:

```
1 include("files/code/figures.jl")
2 pyplot()
```

WARNING: replacing module figures.

Out[39]:

Plots.PyPlotBackend()

Example: Logistic regression


Binary variable, $Y \in \{0, 1\}$, explanatory variable, $X \in \mathbb{R}$.

Example: how is eye strain (Y) related to the number of hours per week (X) spent on MS Teams?

In [40]:

```
1 figures.plot_logistic_regression_data
```

Out[40]:



Assume $Y = \{Y_i, i = 1 \dots n\}$ are id Bernoulli with parameters, p_i , depending on $X = \{X_i, i = 1 \dots n\}$:

$$\mathbb{P}(Y|X) = \prod_{i=1}^n p(X_i)^{Y_i} (1 - p(X_i))^{1-Y_i}$$

Example: Logistic regression

How should Bernoulli parameter, p , vary with X ?

Assume *log odds* are a linear function of X :

$$\log\left(\frac{p(X)}{1 - p(X)}\right) = \beta_0 + \beta_1 X = \beta \cdot \mathbf{X}$$

where $\mathbf{X} = (1, X)$ and $\beta = (\beta_0, \beta_1)$.

Solve for $p(X)$:

$$p(X) = \sigma(\beta \cdot \mathbf{X}),$$

where $\sigma(x)$ is the sigmoid function:

$$\sigma(x) = \frac{1}{1 + \exp(-x)}.$$

Example: Logistic regression as an optimisation problem

Obviously cannot use least squares to select optimal values of β .

Instead maximise, $\mathbb{P}(Y|X)$ - *likelihood of Y given X* - with respect to β .

In fact easier to minimise $-\log \mathbb{P}(Y|X)$ because it turns task into a sum minimisation:

$$\beta_* = \arg \min_{\beta} \sum_{i=1}^n -Y_i \log(\sigma(\beta \cdot \mathbf{X}_i)) - (1 - Y_i) \log(1 - \sigma(\beta \cdot \mathbf{X}_i)).$$

Example: Logistic regression via gradient descent

We can solve this sum minimisation problem with GD or SGD. First we need the gradient of the objective function,

$$L(\beta) = \sum_{i=1}^n -Y_i \log(\sigma(\beta \cdot \mathbf{X}_i)) - (1 - Y_i) \log(1 - \sigma(\beta \cdot \mathbf{X}_i)).$$

Direct calculation gives (check)

$$(\nabla L(\beta))_k = \sum_{i=1}^n (\sigma(\beta \cdot \mathbf{X}_i) - Y_i) (\mathbf{X}_i)_k$$

where $k \in \{0, 1\}$ and $(\mathbf{X}_i)_k$ is the k^{th} component of \mathbf{X}_i .

Implementation: generate some test data

In [41]:

```

1  n = 1000 # Number of data points
2  xrange = [-20.0, 20.0] # Range of values for x
3   $\beta_0$  = [1.0, 1.0] # True parameter values
4
5  sigmoid(x) = 1.0/(1.0+exp(-x))
6
7  # Generate some data
8  Random.seed!(2)
9  X = zeros(n,2)
10 X[:,1] .= 1.0
11 X[:,2] = rand(Uniform(xrange[1], xrange[2]), n)
12 Y = [rand(Bernoulli(sigmoid( $\beta_0$ ·X[i,:])))) for i in 1:n]

```

Out[41]:

1000-element Array{Bool,1}:

```

0
1
0
1
1
1
1
1
0
1
1
0
1
0
:
0
1
0
1
0
0
1
1
0
1
0
1

```

Implementation: define the objective function

Need negative log likelihood of the data (\mathbf{X} , \mathbf{Y}) for a given value of β .

First calculate the components of the sum as separate lists and then assemble them at end:

In [42]:

```

1 function L( $\beta$ )
2     S = [ sigmoid( $\beta \cdot X[i,:]$ ) for i in 1:n]
3     # It *can* happen that S[i]==0.0 or 1.0 due to finite precision. This needs
4     # be handled separately to avoid log(0). Use the ? : ternary operator
5     A1 = [S[i] == 0.0 ? 0.0 : Y[i]*log(S[i]) for i in 1:n]
6     A2 = [S[i] == 1.0 ? 0.0 : (1-Y[i])*log(1.0-S[i]) for i in 1:n]
7     A = [A1[i] + A2[i] for i in 1:n]
8     return -1.0*sum(A)
9 end

```

Out[42]:

L (generic function with 1 method)

Plot of $L(\beta)$


In [43]:

```

1 p = figures.contour_plot(L, [-5.0,5.0], [0.5, 10.0])
2 scatter!([ $\beta_0[1]$ ],[ $\beta_0[2]$ ], marker=:x, markersize=8, label=" $\beta_0$ ", color=:red)

```

Out[43]:



Implementation: define the gradient of the objective function

In [44]:

```

1 # Define the gradient of the likelihood function with respect to  $\beta$ 
2 function dL( $\beta$ )
3     S = [ sigmoid( $\beta \cdot X[i,:]$ ) for i in 1:n]
4     A1 = [(S[i] - Y[i])*X[i,1] for i in 1:n]
5     A2 = [(S[i] - Y[i])*X[i,2] for i in 1:n]
6     return [sum(A1), sum(A2)]
7 end

```

Out[44]:

dL (generic function with 2 methods)

Implementation: define the partial gradient (for SGD)

In [45]:

```

1 # Define the  $i$ th partial gradient of the likelihood function with respect to  $\beta$ 
2 function dL( $\beta, i$ )
3     S = sigmoid( $\beta \cdot X[i,:]$ )
4     A1 = (S - Y[i])*X[i,1]
5     A2 = (S - Y[i])*X[i,2]
6     return [A1, A2]
7 end
8

```

Out[45]:

dL (generic function with 2 methods)

Run gradient descent

This time we will use a pre-defined step size, $\eta_1(k)$, rather than doing full line minimisations.

In [46]:

```

1 include("files/code/GD.jl")
2  $\eta_1(k)$  = 1.0/(1.0 + k/2.0)
3 xpos, ypos = GD.gradient_descent(L, dL, [-5.0, 0.0], 1.0E-6, 200,  $\eta_1$ )
4 println("GD:  $\beta$  = (", xpos[end], ", ", ypos[end], ")." )

```

GD: β = (1.1777019786001126, 0.8687144958918414).

WARNING: replacing module GD.

Run stochastic gradient descent

In [47]:

```

1  $\eta_2(k)$  = 1.0/(1.0 + k/50.0)
2 xpos2, ypos2 = GD.stochastic_gradient_descent(L, dL, [-5.0, 0.0], 500,  $\eta_2$ , n, 1)
3 println("SGD:  $\beta$  = (", xpos2[end], ", ", ypos2[end], ")." )

```

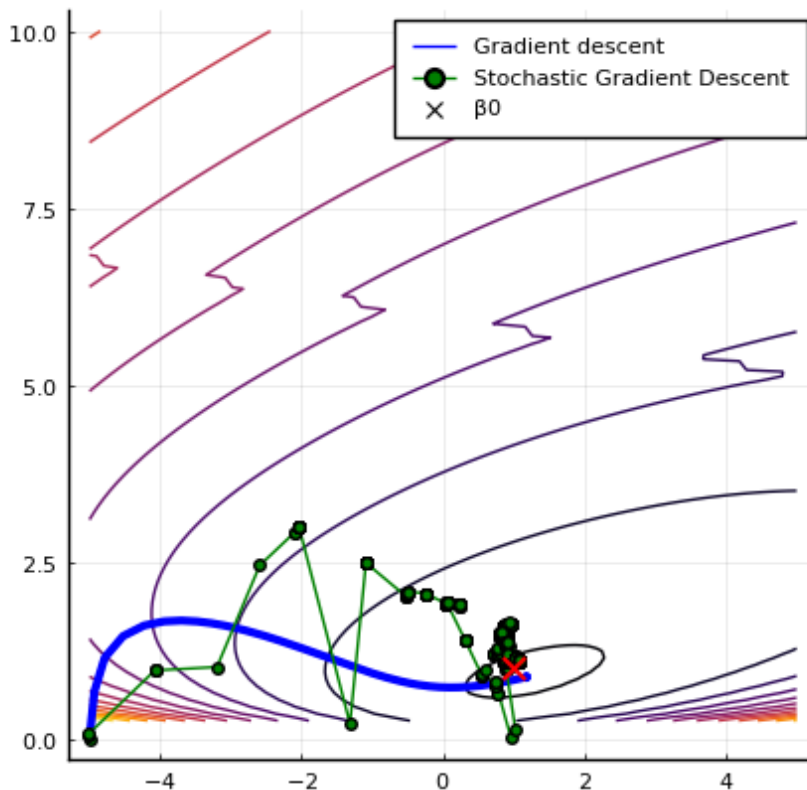
SGD: β = (1.0691395293705568, 1.118280854913653).

Comparison of results

In [48]:

```
1 p = figures.contour_plot(L, [-5.0,5.0], [0.25, 10.0])
2 plot!(xpos, ypos, label="Gradient descent", linewidth=3, color=:blue)
3 plot!(xpos2, ypos2, label="Stochastic Gradient Descent", marker=:circle, color=:green)
4 scatter!([β0[1]], [β0[2]], marker=:x, markersize=8, label="β0", color=:red)
```

Out[48]:



Constrained nonlinear optimisation

Given $f(\mathbf{x}) : \mathbb{R}^n \rightarrow \mathbb{R}$, we can combine what we have learned about gradient descent with what we have learned about linear programming to solve constrained nonlinear optimisation problems of the form:

$$\min_{\mathbf{x} \in C} f(\mathbf{x}),$$

where $C \subset \mathbb{R}^n$ is defined by a set of linear inequalities.

Clearly GD alone is insufficient since \mathbf{x}_n could leave the feasible set, C .

Constrained nonlinear optimisation: Frank-Wolfe algorithm

Idea is to linearise $f(\mathbf{x})$ about \mathbf{x}_n at each iteration of GD and solve the associated LP to obtain an \mathbf{x}_{n+1} , which cannot leave C by construction.

Given current point, \mathbf{x}_n ,

$$f(\mathbf{x}) = f(\mathbf{x}_n) + (\mathbf{x} - \mathbf{x}_n) \cdot \nabla f(\mathbf{x}_n) + \mathcal{O}(|\mathbf{x} - \mathbf{x}_n|^2).$$

Neglecting the $\mathcal{O}(|\mathbf{x} - \mathbf{x}_n|^2)$ terms, minimisation of $f(\mathbf{x})$, is now a LP:

$$\mathbf{x}^* = \arg \min_{\mathbf{x} \in C} [\mathbf{x} \cdot \nabla f(\mathbf{x}_n)]$$

Constrained nonlinear optimisation: Frank-Wolfe algorithm

Start from a point, $\mathbf{x}_0 \in C$. At each iteration, n :

1. Solve the linearised problem

$$\mathbf{x}^* = \arg \min_{\mathbf{x} \in C} [\mathbf{x} \cdot \nabla f(\mathbf{x}_n)]$$

2. Perform *bounded* line minimisation of $f(\mathbf{x})$ from \mathbf{x}_n in the direction of $\mathbf{x}^* - \mathbf{x}_n$:

$$\lambda^* = \arg \min_{\lambda \in [0,1]} f(\mathbf{x}_n + \lambda (\mathbf{x}^* - \mathbf{x}_n))$$

3. Update

$$\mathbf{x}_{n+1} = \mathbf{x}_n + \lambda^* (\mathbf{x}^* - \mathbf{x}_n).$$

4. Repeat until desired tolerance is reached.